

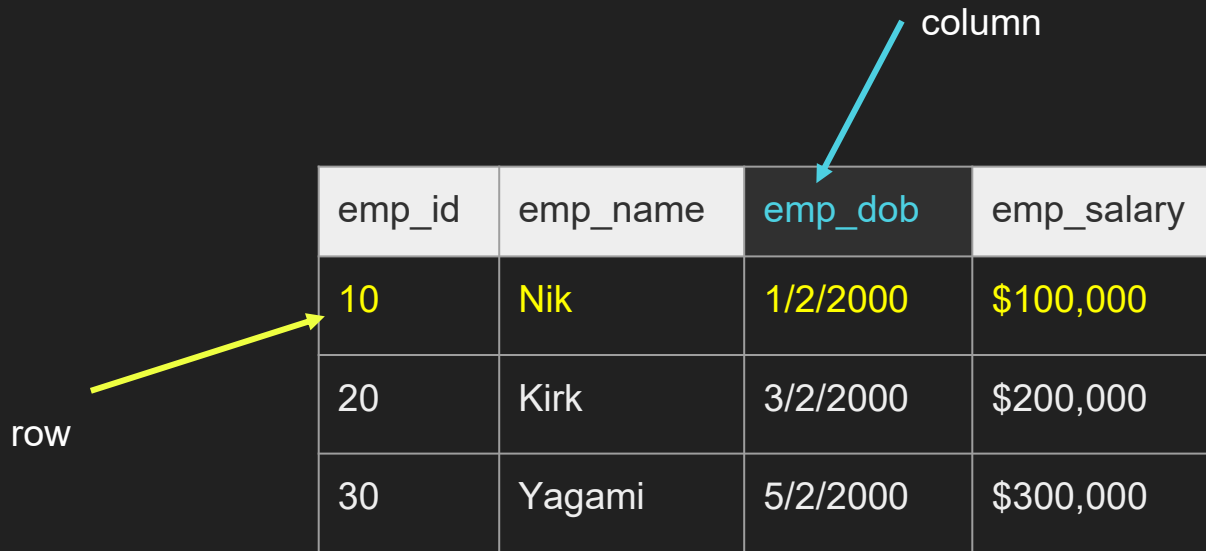
How tables and indexes are stored on disk

And how they are queried

Storage concepts

- Table
- Row_id
- Page
- IO
- Heap data structure
- Index data structure b-tree
- Example of a query

Logical Table



A diagram illustrating a logical table structure. The table has four columns: emp_id, emp_name, emp_dob, and emp_salary. The first row of data is highlighted in yellow. A yellow arrow points to the first row, labeled 'row'. A blue arrow points to the emp_dob column, labeled 'column'.

emp_id	emp_name	emp_dob	emp_salary
10	Nik	1/2/2000	\$100,000
20	Kirk	3/2/2000	\$200,000
30	Yagami	5/2/2000	\$300,000

Row_ID

- Internal and system maintained
- In certain databases (mysql -innoDB) it is the same as the primary key but other databases like Postgres have a system column row_id (tuple_id)

row_id	emp_id	emp_name	emp_dob	emp_salary
1	2000	Nik	1/2/2000	\$100,000
2	3000	Kirk	3/2/2000	\$200,000
3	4000	Yagami	5/2/2000	\$300,000

Page

- Depending on the storage model (row vs column store), the rows are stored and read in logical pages.
- The database doesn't read a single row, it reads a page or more in a single IO and we get a lot of rows in that IO.
- Each page has a size (e.g. 8KB in postgres, 16KB in MySQL)
- Assume each page holds 3 rows in this example, with 1001 rows you will have $1001/3 = 333\sim$ pages

row_id	emp_id	emp_name	emp_dob	emp_salary
1	10	Nik	1/2/2000	\$100,000
2	20	Kirk	3/2/2000	\$200,000
3	30	Yagami	5/2/2000	\$300,000
...
1000	10000	Eddard	1/27/2000	\$250,000

Page 0

1,10,Nik,1/2/2000,
\$100,000|2,
20,Kirk,3/2/2000|3,
30,Yagami,5/2/200
0,\$300,000

Page 1

(Rows 4,5,6)

Page 2

(Rows 7,8,9)

.....

Page 333

More
rows....1000,1000
0,Eddard,1/27/200
0,\$250,000

IO

- IO operation (input/output) is a read request to the disk
- We try to minimize this as much as possible
- An IO can fetch 1 page or more depending on the disk partitions and other factors
- An IO cannot read a single row, its a page with many rows in them, you get them for free.
- You want to minimize the number of IOs as they are expensive.
- Some IOs in operating systems goes to the operating system cache and not disk

Page 0

1,10,Nik,1/2/2000,
\$100,000|2,
20,Kirk,3/2/2000|3,
30,Yagami,5/2/200
0,\$300,000

Page 1

(Rows 4,5,6)

Page 2

(Rows 7,8,9)

.....

Page 333

More
rows....1000,1000
0,Eddard,1/27/200
0,\$250,000

Heap

- The Heap is data structure where the table is stored with all its pages one after another.
- This is where the actual data is stored including everything
- Traversing the heap is expensive as we need to read so much data to find what we want
- That is why we need indexes that help tell us exactly what part of the heap we need to read. What page(s) of the heap we need to pull

Heap

Page 0

1,10,Nik,1/2/1988,
\$100,000|2,
20,Kirk,3/2/1977|3,
30,Yagami,5/2/198
2,\$300,000

Page 1

(Rows 4,5,6)

Page 2

(Rows 7,8,9)

.....

Page 333

More
rows....1000,1000
0,Eddard,1/27/199
9,\$250,000

Index

- An index is another data structure separate from the heap that has “pointers” to the heap
- It has part of the data and used to quickly search for something
- You can index on one column or more.
- Once you find a value of the index, you go to the heap to fetch more information where everything is there
- Index tells you EXACTLY which page to fetch in the heap instead of taking the hit to scan every page in the heap
- The index is also stored as pages and cost IO to pull the entries of the index.
- The smaller the index, the more it can fit in memory the faster the search
- Popular data structure for index is b-trees.

Index on
EMP_ID

IO1 on
the index
to find the
page/row

Page 0

10 (1,0) | 20 (2,0) | 30 (3,0)
40 (4,1) | 50 (5,1) | 60 (6,1)
70 (7,2) | 80 (8,2) | 90 (9,2)

Page 1

100 (10,3) | 110 (11,3) | 120 (12,3)
130 (13,4) | 140 (14,4) | 150 (15,4)
160 (16,5) | 170 (17,5) | 180 (18,5)

.....

Page N

9920 (992,331) | 9930 (993,331) | 9940 (994,331)
9950 (995,332) | 9960 (996,332) | 9970 (997,332)
9980 (998,333) | 9990 (999,333) | 10000 (1000,333)

Heap

IO2 on
the heap
to pull
exactly
the
page(s)
we found
in the
index

Page 0

1,10,Nik,1/2/2000,
\$100,000|2,
20,Kirk,3/2/2000|3,
30,Yagami,5/2/200
0,\$300,000

Page 1

(Rows 4,5,6)

Page 2


(Rows 7,8,9)

.....

Page 333

More
rows....1000,1000
0,Eddard,1/27/200
0,\$250,000

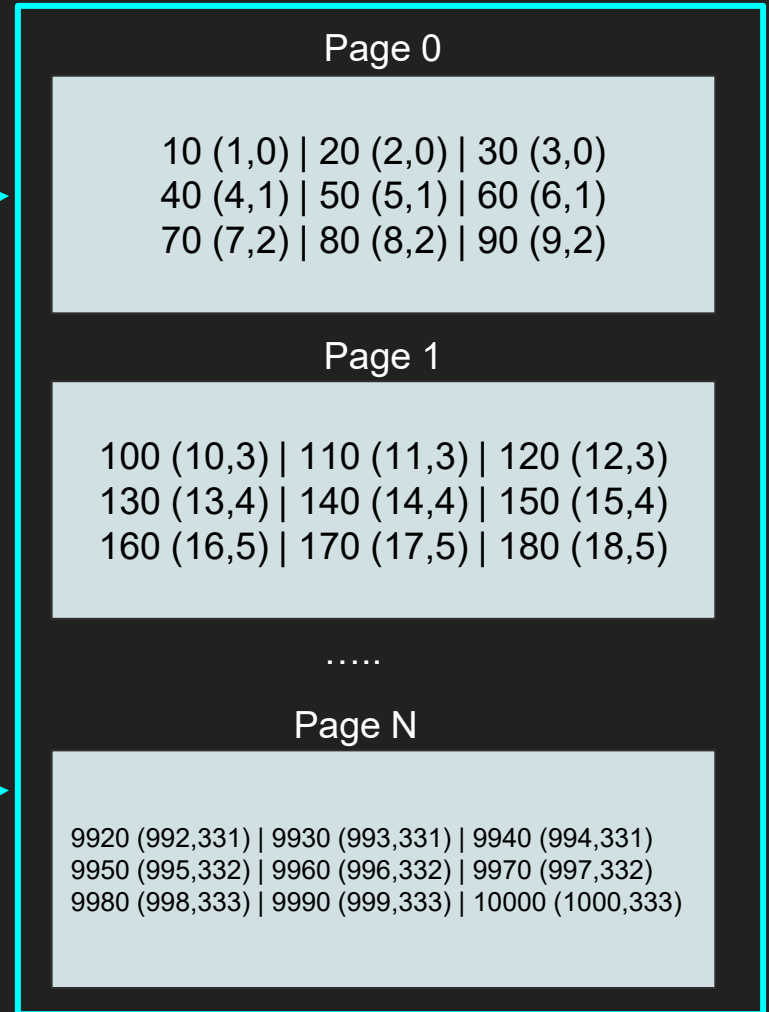
No Index -
SELECT * FROM EMP
WHERE EMP_ID =
10000;



With Index -
`SELECT * FROM EMP
WHERE EMP_ID =
10000;`

10000 (1000,333)

Index on
EMP_ID



10000 (1000,333)

Fetch page 333, and pull row
10000

With Index -

```
SELECT * FROM EMP  
WHERE EMP_ID =  
10000;
```

Heap

Page 0

1,10,Nik,1/2/1988,
\$100,000|2,
20,Kirk,3/2/1977|3,
30,Yagami,5/2/198
2,\$300,000

Page 1

(Rows 4,5,6)

Page 2

(Rows 7,8,9)

.....

Page 333

More
rows....1000,1000
0,Eddard,1/27/199
9,\$250,000



Notes

- Sometimes the heap table can be organized around a single index. This is called a clustered index or an Index Organized Table.
- Primary key is usually a clustered index unless otherwise specified.
- MySQL InnoDB always have a primary key (clustered index) other indexes point to the primary key “value”
- Postgres only have secondary indexes and all indexes point directly to the row_id which lives in the heap.

Storage concepts - Summary

- Table
- Row_id
- Page
- IO
- Heap data structure
- Index data structure b-tree
- Example of a query